

Sumit Kumar

📍 Bengaluru, India ✉ sklegacy789@gmail.com ☎ +91-7033241380 📱 in/sumitkumar22 🌐 sumit22012004.github.io/Portfolio/

SUMMARY

Full-Stack AI Engineer delivering end-to-end intelligent systems that combine automation, reasoning, and adaptive decision-making. I transform complex business needs into scalable AI products, improve operational efficiency, and drive measurable outcomes through reliable architecture, autonomous workflows, and continuous optimization aligned with organizational goals.

EXPERIENCE

Senior AI Engineer (Freelance)

India Career Support | Remote

December 2025 – Present

- Architected and productionized a scalable career-counseling and job-support chatbot using fine-tuned LLMs (LoRA/QLoRA) for domain adaptation, automating resume screening, interview simulation, and personalized career plans — reduced manual counselor intervention by ~45% and handled thousands of monthly user sessions.
- Engineered a hybrid RAG retrieval layer combining Qdrant vector search with MongoDB metadata stores, and exposed retrieval/orchestration via FastAPI microservices; applied systematic context-engineering and template-based prompt strategies to increase answer relevance and reduce hallucinations by ~30%.
- Orchestrated multi-agent workflows (Skills-Assessor, Job-Matcher, Interview-Simulator) using LangGraph and LangChain, deployed the solution on Azure (AKS, Functions, Blob Storage) with CI/CD and Application Insights monitoring, achieving autoscaled inference and sub-second average response times.

AI Engineer

AIRA | Bangalore, India

September 2024 – November 2025

- Engineered a privacy-preserving LLM infrastructure integrating encrypted memory graphs, vector retrieval, and real-time context reconstruction, enabling long-duration conversations with a >92% continuity score across multi-session interactions.
- Developed a modular orchestration layer that unified GPT-5, Gemini, and offline on-prem LLMs through adaptive routing, compression-aware token budgeting, and latency-optimized retrieval, resulting in a 38% improvement in inference throughput.
- Implemented autonomous agent coordination using custom state machines and event-driven pipelines, enabling self-repair routines, memory consistency checks, and automated data-sanitization cycles, reducing manual maintenance efforts by 60%.

SKILLS

AI Engineering : Python, LLM Fine-Tuning, Model Optimization, Prompt Engineering, Context Engineering, Multi-Agent Systems, RAG Pipelines, Vector Search, Semantic Retrieval, Memory Architectures, Inference Optimization

Frameworks & Libraries : LangChain, LangGraph, TensorFlow, PyTorch, Hugging Face, Scikit-Learn, OpenCV, NumPy, Pandas, Matplotlib

Backend & API Development : FastAPI, RESTful API Design, Microservices Architecture, Authentication & Authorization, Asynchronous Processing, Testing & Debugging

Data Engineering : Data Preprocessing, Feature Engineering, ETL Pipelines, Data Cleaning, Model Evaluation, Web Scraping

Databases & Storage : Qdrant, Pinecone, MongoDB, MySQL, Redis, Vector Databases, Document Stores, Metadata Indexing

Cloud & DevOps : Azure, AWS, CI/CD Pipelines, Docker, Deployment Automation, Logging & Monitoring, API Performance Optimization, Git/GitHub

Professional Competencies : Analytical Thinking, Problem Solving, System Design Mindset, Product Ownership, Team Collaboration, Communication

PROJECT

Intelligent Multi-Agent Research Assistant

- Designed a cooperative swarm of role-specialized agents capable of autonomous document ingestion, semantic clustering, and progressive summarization using Qdrant + Pinecone retrieval layers, reducing research review cycles by 40%.
- Built a persistent contextual memory fabric using LangGraph pipelines that allowed cross-task continuity, reinforcement-based evidence selection, and historical query referencing, achieving 95% semantic relevance across multi-turn evaluations.
- Delivered an interactive research environment with a FastAPI backend and Streamlit UI, integrating real-time insight graphs, citation extraction, and iterative refinement loops to significantly accelerate domain-specific knowledge synthesis.

EDUCATION

Bachelor of Engineering – Computer Science (AI & ML Specialization)

Chandigarh University | 2021 – 2025